


Problem

- Manual moderation does not scale
- High latency in reacting to toxic content
- No visibility into model behavior



Comments Toxicity Detection

Solution

- Classifies comments as TOXIC / CLEAN
- Returns toxicity probability
- Asynchronous processing via Kafka
- Persistent storage and monitoring

Demo: docker compose up -d | UI: localhost:8501 |
Grafana: localhost:3000



Tech Stack

